

# Przetwarzanie dużych grafów za pomocą Apache Giraph

Marek Rogala

Wrzesień 2014

## **Streszczenie**

## **Słowa kluczowe**

graph queries, Datalog

## **Dziedzina pracy (kody wg programu Socrates-Erasmus)**

11.3 Informatyka

## **Klasyfikacja tematyczna**

D. Software

D.127. Blabalgorithms

D.127.6. Numerical blabalysis

## **Tytuł pracy w języku angielskim**

Processing large graphs with Apache Giraph



# Spis treści

<b>Introduction</b> . . . . .	5
<b>1. Datalog</b> . . . . .	7
1.1. Definitions . . . . .	7
1.1.1. Queries and instances . . . . .	7
1.1.2. Datalog with negation . . . . .	7
1.1.3. Positive and semi-positive Datalog . . . . .	8
1.1.4. Stratified semantics . . . . .	8
1.2. Evaluation strategies for Datalog . . . . .	9
<b>2. Pregel/Giraph</b> . . . . .	11
<b>3. Socialite</b> . . . . .	13
3.1. Motivation . . . . .	13
3.2. Meet operation and induced ordering . . . . .	14
3.2.1. Order induced by a meet operation . . . . .	14
3.2.2. Order on relation instances . . . . .	14
3.3. Socialite program . . . . .	15
3.4. Semantics and evaluation - one relation case . . . . .	15
<b>4. Translating Socialite programs to Giraph programs</b> . . . . .	17
<b>5. Implementation</b> . . . . .	19
<b>6. Summary</b> . . . . .	21
<b>Bibliografia</b> . . . . .	23



# Introduction

In recent years, the humanity has created many graph datasets much larger than those available ever before. Those graphs became a very popular object of research. Most notable examples include *the Web graph* – a graph of internet websites and links between them, all kinds of social networks. There are also graphs such as transportation routes, similarity of newspaper of scientific articles or citations among them.

With increasing computational power and memory space, we can expect more and more real-life graphs to become subject to computation. We can also expect the existing graphs, such as the Web or social networks, to grow in all aspects.

The graphs mentioned can be a source of a huge amount of useful information. Hence, there is an increasing number of practical computational problems. Some of the analyses carried out are ranking of the graph nodes (ex. importance of a Web page, determining most influential users in given group of people), clustering (detecting communities), computing metrics for the whole graph or some parts of it (ex. connectivity measures) and connection predictions. Usually, such analyses are built on top of standard graph algorithms, such as PageRank-like procedures, shortest paths or connected components.

In the past, we have seen many tools for efficient distributed large dataset computations, starting from Google's MapReduce [?] and its widely used open source counterpart, Apache's Hadoop, as well as higher-level languages such as PigLatin and Hive. However, those are not well suited for graph computations, as they do not support iteration well.

Recently, there is an outbreak of frameworks and languages for large graphs processing, including industrial systems such as Google's Pregel [?], its open-source version Apache Giraph and Spark GraphX, Giraph++ [?].

The frameworks currently available allow you to implement a graph algorithm in a specified model (for example Pregel's "think like a vertex"). On the other hand, query languages, such as SQL, are a bad fit for graph data since because of limited support of iteration. With the rise of graph computational problems, we need an easier way to extract information from graphs: a query language for effectively expressing data queries typical for graphs.

The Socialite [SL] [DSL] language is one of the most interesting propositions. It is based on a classical language – Datalog. Declarative semantics makes it easy to distribute the computations, since no execution flow is embedded in the program code. It also gives many possibilities for optimizations and approximate evaluation. At the same time, support for recursion is crucial, since most graph algorithms have iterative nature. However, most of practical graph algorithms can not be expressed efficiently in Datalog because of the language limitations. With a few extensions to original Datalog, most important of which is recursive aggregation, Socialite makes it easy to write intuitive programs which can be executed very efficiently.

Unfortunately, there is no solid implementation of Socialite available. The interpreter published by the authors is undocumented and contains many bugs. It is hard to imagine the language being implemented in the industry in the foreseeable future. The papers [SL] and

[DSL] which introduced Socialite contain certain simplifications and are not specific about some important details in definitions and proofs.

The goal of this work is to bridge the gap between the theoretical idea for Socialite and a practical implementation and to draw a path for its usage in the industry. We show how to translate Socialite declarative programs into Giraph "think-like-a-vertex" programs and introduce a compiler that enables Socialite programs to be executed on existing Giraph infrastructure. This allows users of Hadoop to write and execute Socialite programs without any additional effort to build a dedicated server infrastructure for that.

The work consists of six chapters. In 1 we recall definitions of Datalog and its evaluation methods while 2 contains an introduction to the Pregel computation model. In 3 we describe the extensions introduced by Socialite and provide formal definitions and general-case proofs which the original papers lack. Chapter 4 shows the translation procedure from Socialite to Giraph programs implemented in the S2G compiler, which is described in 5. In 6 we sketch the possible future work and the path to industrial implementation of the language using the S2G compiler.

Socialite Implementation Summary

# Rozdział 1

## Datalog

In this chapter we define the Datalog language.

### 1.1. Definitions

@TODO: te definicje są roboczo przepisane z [WFoM], którzy z kolei podają je za [FoD]].

#### 1.1.1. Queries and instances

Let us assume that **dom** is an infinite set of data values.

- A *database schema*  $\sigma$  is a collection of relation names  $R$  where every  $R$  has arity  $ar(R) \geq 0$ .
- We call  $R(\bar{d})$  a *fact* when  $R$  is a relation name and  $\bar{d}$  is a tuple in **dom**.
- We say that a fact  $R(d_1, \dots, d_k)$  is *over* a database schema  $\sigma$  if  $R \in \sigma$  and  $ar(R) = k$ .
- A (*database*) *instance*  $I$  over  $\sigma$  is a finite set of facts over  $\sigma$ .
- We denote by  $adom(I)$  the set of all values that occur in facts of  $I$ . When  $I = \mathbf{f}$ , we simply write  $adom(\mathbf{f})$  rather than  $adom(\{\mathbf{f}\})$ .
- By  $|I|$  we denote the number of facts in  $I$ .

#### 1.1.2. Datalog with negation

Let **var** be the universe of variables, disjoint from **dom**. An atom is of the form  $R(u_1, \dots, u_k)$  where  $R$  is a relation name and each  $u_i \in \mathbf{var}$ . We call  $R$  the predicate. A *literal* is an atom (*positive atom*) or a negated atom (*negative atom*).

We recall Datalog with negation [FoD], abbreviated Datalog<sup>−</sup>. Formally, a Datalog<sup>−</sup> rule  $\phi$  is a quadruple  $(head_\phi, pos_\phi, neg_\phi, ineq_\phi)$  where  $head_\phi$  is an atom;  $pos_\phi$  and  $neg_\phi$  are sets of atoms;  $ineq_\phi$  is a set of inequalities  $(u = v)$  with  $u, v \in \mathbf{var}$  and the variables of  $\phi$  all occur in  $pos_\phi$ . The components  $head_\phi$ ,  $pos_\phi$  and  $neg_\phi$  are called respectively the *head*, the *positive body atoms* and the *negative body atoms*. We refer to  $pos_\phi \cup neg_\phi$  as the *body atoms*. Note,  $neg_\phi$  contains just atoms, not negative literals. Every Datalog<sup>−</sup> rule  $\phi$  must have a head,  $pos_\phi$  must be non-empty and  $neg_\phi$  may be empty. If  $neg_\phi = \emptyset$  then  $\phi$  is called positive.



Of course, a rule  $\phi$  may be written in the conventional syntax. For instance, if  $head_\phi = T(u, v)$ ,  $pos_\phi = \{R(u, v)\}$ ,  $neg_\phi = \{S(v)\}$ , and  $ineq_\phi = u \neq v$ , with  $u, v \in var$ , then we can write  $\phi$  as:

$$T(u, v) : \neg R(u, v), \neg S(v), u \neq v$$

The set of variables of  $\phi$  is denoted  $vars(\phi)$ . A rule  $\phi$  is said to be over schema  $\sigma$  if for each atom  $R(u_1, \dots, u_k) \in head_\phi \cup pos_\phi \cup neg_\phi$ , the arity of  $R$  in  $\phi$  is  $k$ . A Datalog program  $P$  over  $\sigma$  is a set of Datalog eg rules over  $\sigma$ . We write  $sch(P)$  to denote the (minimal) database schema that  $P$  is over. We define  $idb(P) \subset sch(P)$  to be the database schema consisting of all relations in rule-heads of  $P$ . We abbreviate  $edb(P) = sch(P) \setminus idb(P)$ . As usual, the abbreviation "idb" stands for "intensional database schema" and "edb" stands for "extensional database schema". A valuation for a rule  $\phi$  in  $P$  w.r.t. an instance  $I$  over  $edb(P)$ , is a total function  $V : vars(\phi) \rightarrow dom$ . The application of  $V$  to an atom  $R(u_1, \dots, u_k)$  of  $\phi$ , denoted  $V(R(u_1, \dots, u_k))$ , results in the fact  $R(a_1, \dots, a_k)$  where  $a_i = V(u_i)$  for each  $i \in 1, \dots, k$ . This is naturally extended to a set of atoms, which results in a set of facts. The valuation  $V$  is said to be satisfying for  $\phi$  on  $I$  if  $V(pos_\phi) \subset I$ ,  $V(neg_\phi) \cap I = \emptyset$ , and  $V(u) \neq V(v)$  for each  $(u \neq v) \in ineq_\phi$ . If so,  $\phi$  is said to derive the fact  $V(head_\phi)$ .

### 1.1.3. Positive and semi-positive Datalog

A Datalog  $\neg$  program  $P$  is positive if all rules of  $P$  are positive. We say that  $P$  is semi-positive if for each rule  $\phi \in P$ , the atoms of  $neg_\phi$  are over  $edb(P)$ . We now give the semantics of a semi-positive Datalog  $\neg$  program  $P$ . First, let  $T_P$  be the immediate consequence operator that maps each instance  $J$  over  $sch(P)$  to the instance  $J = J \cup A$  where  $A$  is the set of facts derived by all possible satisfying valuations for the rules of  $P$  on  $J$ . Let  $I$  be an instance over  $edb(P)$ . Consider the infinite sequence  $I_0, I_1, I_2$ , etc, inductively defined as follows:  $I_0 = I$  and  $I_i = T_P(I_{i-1})$  for each  $i \geq 1$ . The output of  $P$  on input  $I$ , denoted  $P(I)$ , is defined as  $\bigcup_j I_j$ ; this is the *minimal fixpoint* of the  $T_P$  operator.

### 1.1.4. Stratified semantics

We say that  $P$  is syntactically stratifiable if there is a function  $\rho : sch(P) \rightarrow 1, \dots, |idb(P)|$  such that for each rule  $\phi \in P$ , having some head predicate  $T$ , the following conditions are satisfied:

1.  $\rho(R) \leq \rho(T)$  for each  $R(\bar{u}) \in pos_\phi \cap idb(P)$
2.  $\rho(R) < \rho(T)$  for each  $R(\bar{u}) \in neg_\phi \cap idb(P)$ .

For  $R \in idb(P)$ , we call  $\rho(R)$  the stratum number of  $R$ . Intuitively,  $\rho$  partitions  $P$  into a sequence of semi-positive Datalog  $\neg$  programs  $P_1, \dots, P_k$  with  $k \leq |idb(P)|$  such that for each  $i = 1, \dots, k$ , the program  $P_i$  contains the rules of  $P$  whose head predicate has stratum number  $i$ . This sequence is called a syntactic stratification of  $P$ . We can now apply the stratified semantics to  $P$ : for an input  $I$  over  $sch(P)$ , we first compute the fixpoint  $P_1(I)$ , then the fixpoint  $P_2(P_1(I))$ , etc. The output of  $P$  on input  $I$ , denoted  $P(I)$ , is defined as  $P_k(P_{k-1}(\dots P_1(I) \dots))$ . It is well known that the output of  $P$  does not depend on the chosen syntactic stratification (if more than one exists). Not all Datalog  $\neg$  programs are syntactically stratifiable. By *stratified Datalog* we refer to all Datalog  $\neg$  programs which are syntactically stratifiable.

## 1.2. Evaluation strategies for Datalog

bottom-up, top-down



## Rozdział 2

# Pregel/Giraph

@TODO: Description of Pregel model and info about the Giraph project, also mention other extensions like Giraph++.



## Rozdział 3

# Socialite

Socialite ([SL], [DSL]) is a graph query language based on Datalog. While Datalog allows to express some of graph algorithms in an elegant and succinct way, many practical problems cannot be efficiently solved with Datalog programs. Socialite consists of a few extensions to the Datalog language, most significant of which is the ability to combine recursive rules with aggregation. Under some conditions, such rules can be evaluated incrementally and thus as efficiently as regular recursion in Datalog.

### 3.1. Motivation

Most graph algorithms are essentially some kind of iteration or recursive computation. In many of them, the computation results are gradually refined in each iteration, until the final result is reached. Examples of such algorithms are the Dijkstra algorithm for single source shortest paths or Page Rank. An important limitation of Datalog is the lack of possibility to create recursive aggregation rules.

A straightforward Datalog program for computing single source shortest paths (starting from 0) is given below. Due to limitations of Datalog, this program is faulty: it computes all possible path lengths from node 1 (the source) to other nodes in the first place, and after that for each node the minimal distance is chosen. Not only this approach results in bad performance, but causes the program to execute infinitely if a loop is reachable from the source.

$\text{PATH}(t, d)$	: –	$\text{EDGE}(1, t, d).$
$\text{PATH}(t, d)$	: –	$\text{PATH}(s, d_1), \text{EDGE}(s, t, d_2), d = d_1 + d_2.$
$\text{MINPATH}(t, \text{MIN}(d))$	: –	$\text{PATH}(t, d).$

Rysunek 3.1: Datalog query for computing shortest paths from node 1 to other nodes

Socialite allows aggregation to be combined with recursion under some conditions. This allows us to write straightforward programs for such problems, which finish execution in finite time and often are much more efficient than Datalog programs. An example Socialite program that computes single source shortest paths is presented below.

While being very useful, recursive aggregation rules have a unambiguous solutions only under some conditions on the rules and the aggregation function itself.

PATH( $t, \text{MIN}(d)$ )	: −	EDGE( $1, t, d$ ).
	: −	PATH( $s, d_1$ ), EDGE( $s, t, d_2$ ), $d = d_1 + d_2$ .

Rysunek 3.2: Socialite query for computing shortest paths from node 1 to other nodes

Typically, Datalog programs semantics is defined using the least fixed point of instance inclusion. This requires that the subsequent computation iterations only add tuples to the database instance, but never remove tuples from the instance. When recursive aggregate functions are allowed, this is not the case: a tuple in the instance can be replaced with a different one because a new aggregated value appeared. Consequently, in order to define Socialite programs semantics in terms of least fixed point, we need to use a different order on database instances.

First, we will define a meet operation and show the order that it induces. Then we will show that if the aggregation function is a meet operation and corresponding rules are monotone with respect to this induced order, then the result of the program is unambiguously defined. We will also show how it can be computed efficiently.

## 3.2. Meet operation and induced ordering

**Definition 3.2.1.** A binary operation is a *meet* operation if it is idempotent, commutative and associative.

@TODO: Maybe remind definitions of semi-lattice and partial order?

### 3.2.1. Order induced by a meet operation

A meet operation  $\sqcap$  defines a semi-lattice: it induces a partial order  $\preceq$  over its domain, such that the result of the operation for any two elements is the least upper bound of those elements with respect to  $\preceq$

**Example 3.2.1.**  $\max(a, b)$  for  $a, b \in \mathbb{N}$  is a meet operation; it is:

- idempotent –  $\max(a, a) = a$
- commutative –  $\max(a, b) = \max(b, a)$
- associative –  $\max(a, \max(b, c)) = \max(\max(a, b), c)$

It induces the partial order  $\leq$ : for any two  $a, b \in \mathbb{N}$ ,  $\max(a, b)$  is their least upper bound with respect to  $\leq$ .

On the contrary,  $+$  is not a meet operation, since it is not idempotent:  $1 + 1 \neq 1$ .

### 3.2.2. Order on relation instances

We can extend the partial order  $\preceq$  to an order  $\sqsubseteq$  on relation instances as follows:

$$R_1 \sqsubseteq R_2 \iff \forall_{(q_1, \dots, q_{n-1}, v) \in g(R_1)} \exists_{(q_1, \dots, q_{n-1}, v') \in g(R_2)} v \preceq v'$$

@TODO: Show that this is actually a partial order.

@TODO: komentarz, przykład

We can easily see that an empty relation instance  $\emptyset$  is smaller under  $\sqsubseteq$

### 3.3. Socialite program

A Socialite program is a Datalog program, with additional aggregation function defined for some of the relations: For each relation  $R$ , there can be one column  $agg_R$  chosen for which an aggregation function  $\sqcap_R$  is provided. The rest of the columns are called the *qualifying columns*. Intuitively, after each step of computation, we group the tuples in the relation by the qualifying columns and aggregate the column  $agg_R$  using  $\sqcap_R$ .

For simplicity, we assume that the aggregated column is always the last one.

Syntactically, we require that all rules with a given relation  $R$  in the head are defined together, with the head defined once and possibly many rule bodies on the right hand side:

$$\begin{array}{lll} P(x_1, \dots, x_k, \text{AGG}(y)) & : - & Q_1(x_1, \dots, x_k, y) \\ & \dots & \\ & : - & Q_m(x_1, \dots, x_k, y). \end{array}$$

Rysunek 3.3: Definition of multiple rules for one relation.

### 3.4. Semantics and evaluation - one relation case

In this section we will show that the semantics of a Socialite program can be unambiguously defined using least fixed point, as long as it satisfies some conditions. To simplify the reasoning, we will restrict our attention to programs with only one *idb* relation. In the next section we extend the definitions and theorems from this section to the general case of many *idb* relations.

Let  $P$  be a Socialite program, with only one *idb* relation  $R$  of arity  $k$  in the form of:

$$\begin{array}{lll} R(x_1, \dots, x_{k-1}, F(y)) & : - & Q_1(x_1, \dots, x_{k-1}, y) \\ & \dots & \\ & : - & Q_m(x_1, \dots, x_{k-1}, y). \end{array}$$

Let us define:

- $f : \mathbf{dom}^k \rightarrow \mathbf{dom}^k$  – the function that evaluates the rules  $Q_1, \dots, Q_m$  based on the contents of the relation given in the argument and returns its input extended with the set of generated tuples
- $g : \mathbf{dom}^k \rightarrow \mathbf{dom}^k$  – the function that performs grouping and aggregation:

$$g(R) = \{(x_1, \dots, x_{k-1}, F(\{y : (x_1, \dots, x_{k-1}, y) \in R\})) : (x_1, \dots, x_{k-1}, \bar{y}) \in R\}$$

- $h = g \circ f$

**Theorem 3.4.1.** *If  $f$  is monotone with respect to  $\sqsubseteq$ , i.e.  $R_1 \subseteq R_2 \rightarrow f(R_1) \subseteq f(R_2)$ , and there exists  $n \geq 0$ , such that  $h^n(\emptyset) = h^{n+1}(\emptyset)$ , then  $R^* = h^n(\emptyset)$  is the least fixed point of  $h$ , that is:*



1.  $R^* = h(R^*)$

2.  $R^* \sqsubseteq R$  for all  $R$  such that  $R = h(R)$

*Proof:*

$g$  is monotone with respect to  $\sqsubseteq$ . Since we assumed that  $f$  is monotone with respect to  $\sqsubseteq$ ,  $h = g \circ f$  is also monotone with respect to  $\sqsubseteq$ .

Let us notice that  $\emptyset$  is greater under  $\sqsubseteq$  than any other element. TODO: show that (very easy)

Let us suppose that  $R'$  is any fixpoint of  $h$ . We know that  $\emptyset \sqsubseteq R'$ . Applying  $h$   $n$  times to both sides of the inequality, we have that  $R^* = h^n(\emptyset) \sqsubseteq h^n(R') = R'$ , thanks to the monotonicity of  $h$  with respect to  $\sqsubseteq$ . Therefore, the inductive fixed point  $R^*$  is the least fixed point of  $h$ .

@TODO: open questions:

- How we compute recursive functions with non-meet aggregation operators?

## Rozdział 4

# Translating Socialite programs to Giraph programs



## Rozdział 5

# Implementation



## Rozdział 6

## Summary



# Bibliografia

- [SL] Jiwon Seo, Stephen Guo, Monica S. Lam, *SociaLite: Datalog extensions for efficient social network analysis*, ICDE 2013: 278-289
- [DSL] Jiwon Seo, Jongsoo Park, Jaeho Shin, Monica S. Lam: *Distributed SociaLite: A Datalog-Based Language for Large-Scale Graph Analysis*. PVLDB 6(14): 1906-1917 (2013)
- [FoD] S. Abiteboul, R. Hull, and V. Vianu: *Foundations of Databases*. Addison-Wesley (1995)
- [WFOm] T.J. Ameloot, B. Ketsman, F. Neven, D. Zinn: *Weaker Forms of Monotonicity for Declarative Networking: a more fine-grained answer to the CALM-conjecture*, PODS (2014) ...???..